

Reducing SQIsign v2 Memory and Transferring Lifetime Scheduling to v3

An Order-Preserving Norm Sketch and RP2350 Evaluation

Hiro Nakanishi

Independent Researcher

quantumsity@protonmail.com

Preprint version: 5 September 2026

Abstract

Small SQIsign keys and signatures do not imply a small signing workspace. In the SQIsign v2.0.1 Level-I compact implementation studied here, five simultaneously live regions in `find_uv` total 6,339,868 bytes, exceeding RP2350 on-chip SRAM. We combine typed lifetime scheduling with an order-preserving norm sketch that stores a bit length and 64-bit prefix and regenerates full precision only when required.

For Level I/RADIX32, the operation arena falls from 353,008 to 172,080 bytes (51.2532%). The firmware uses no GMP, dynamic allocation, or external PSRAM; it reserves 321,408 of 532,480 SRAM bytes and leaves 211,072 unreserved. The ordinary API and two separately started low-memory runs process the same 100 official requests. These are fresh-process repetitions of a differential check, not independent or historical-response oracles. For the fixed image, a linked call-graph analysis bounds the KeyGen, Sign, and Verify operation-PSP paths at 92,192, 120,664, and 20,980 bytes, including one 212-byte exception entry; handler frames, interrupt nesting, and live MSP use are excluded.

For v3.0, we overlay three lifetime-disjoint pairs in two functions. Both variants pass API, self-test, and official-response checks on 300 distinct parameter-set/request pairs each, giving 600 implementation-vector runs in total, and all six affected Arm frames shrink. One clean `p324_3` RP2350 path reduces the Sign PSP watermark from 101,060 to 96,396 bytes. These target data cover one parameter and one deterministic path, not a timing distribution or whole-program stack bound. Software screens find input-dependent behavior; we make no claim of constant-time or power/EM resistance.

Keywords: SQIsign, embedded cryptography, memory optimization, lifetime scheduling, microcontroller, side-channel analysis

1 Introduction

SQIsign is a Round-3 candidate in NIST’s additional post-quantum signature process[1, 2]. It obtains unusually small public keys and signatures from supersingular-isogeny and quaternion computations [3, 4, 5]. This communication compactness is attractive for constrained systems, but the signing path still combines quaternion arithmetic, four-dimensional lattice routines, norm equations, candidate enumeration, and large-integer operations.

Fixed-precision arithmetic removes unbounded integer growth, but replacing a heap object with a maximum-sized stack array does not by itself reduce total SRAM. In our v2 starting point, five live allocations in one `find_uv` invocation total 6,339,868 bytes, about 11.9 times the RP2350’s 532,480-byte SRAM. The central systems question is therefore which values must coexist, not merely which allocator stores them.

This paper studies four questions:

- Q1. Can a complete deterministic SQIsign v2 Level-I KeyGen/Sign/Verify path run in RP2350 on-chip SRAM without GMP, a heap, or external PSRAM?
- Q2. Can the resident table of full-precision candidate norms be removed while preserving the reference comparison and search order?
- Q3. What are the accompanying SRAM, stack, flash, time, and side-channel observations, and which conclusions do they support?
- Q4. Does lifetime scheduling remain useful after the substantial v3 redesign, where the v2 candidate table no longer exists?

Our contributions are as follows.

1. We formulate protocol workspace placement as an aligned storage allocation problem over typed objects and path-sensitive lifetimes. A manually reviewed contract is compiled into C unions, offsets, and static assertions.
2. We introduce an order-preserving norm sketch for the v2 candidate table. It reduces resident state while retaining a full-precision fallback whenever the sketch alone cannot decide an ordering.
3. We produce an auditable v2 artifact: source correspondence checks, lifetime certificates, sanitizer and guard tests, ELF audits, two fresh-process repetitions of a 100-request differential campaign, a fixed-image operation-PSP analysis, and RP2350 captures.
4. We transfer lifetime scheduling, but not the v2 norm sketch, to two v3 functions. We test all three official parameter sets on the host and in Arm object builds, and report one clean p324_3 target path.
5. We separate functionality and memory claims from implementation security. The measured implementations remain variable-time, and no analog trace campaign has been performed.




Figure 1: Study overview. The v2 path combines lifetime scheduling with an order-preserving norm sketch. The v3 D3 experiment transfers only the scheduling principle. Functionality and memory evidence are kept separate from the side-channel claim boundary.

2 Related Work

Reusing storage for noninterfering lifetimes is an established compiler and memory-allocation principle. Graph-coloring register allocation represents simultaneously live values by interference edges[6], and compile-time memory allocation has considered differently sized objects [7]. We do not claim to invent lifetime-based reuse. Our focus is its application to complete SQIsign operation closures, including typed layout, failure paths, clearing, linked-image accounting, and a reconstructible evidence chain.

Fixed-precision SQIsign first supplied uniform bounds of 7026/10713/14150 bits for Levels I/III/V[8]. Compact Quaternion Algorithms later reduced these values to 1665/2521/3319 bits and improved host performance [9]. Qlapoti and follow-up work improve the ideal-to-isogeny translation and its analysis [10, 11, 12]. These advances bound or accelerate arithmetic objects; our v2 contribution instead targets their simultaneous residency across the protocol.

For embedded SQIsign, pqm4 documented the difficulty of integrating the v2 signing side because of GMP and dynamic allocation[13]. Cortex-M work has optimized finite-field arithmetic or one-dimensional verification[14, 15], while a complete KeyGen/Sign/Verify implementation has been reported for larger Arm systems [16]. Other nearby implementations include a 32-bit deployment study[17], an AVX-512 implementation [18], and a pure-Rust v2 implementation that still uses heap-backed integers[19]. The v3 submission itself now includes complete Cortex-M4 paths without external GMP [5, 20]. Accordingly, the v2 on-chip result and the v3 transfer experiment are reported as distinct contributions, not as a performance comparison between generations.

Constant-time integer, lattice, and quaternion components for SQIsign are active research topics[21, 22, 23]. Mukherjee et al. demonstrated a simple-power-analysis key-recovery route via a secret-derived exponent in Cornacchia processing[24]. These works show why a static address layout is not a side-channel countermeasure. Our timing and software-trace screens identify residual variable behavior; they do not replace power/EM acquisition or attack reproduction.

Our literature audit was frozen on 2026-09-04, after the v3 release. Search queries, databases, inclusion criteria, and nearby counterexamples are stored in the public artifact. The negative result is deliberately narrow: we found no prior report matching the combined v2 scope of complete K/S/V, no GMP or heap, on-chip-only RP2350 execution, and linked-image accounting. It is not a claim over unpublished or later work.

3 Background and Claim Boundary

3.1 SQIsign v2 and v3

The v2.0.1 implementation studied here uses a two-dimensional response path, including a chain of (2,2)-isogenies, and is not the earlier one-dimensional verification setting[4]. We fix Level I and the RADIX32 finite-field backend. Quaternion integers use 27 64-bit limbs: a 1665-bit magnitude bound plus sign storage rounded to a 1728-bit array. Multiplication, division, GCD, square root, and modular exponentiation operate on these fixed-capacity arrays and caller-owned workspaces.

SQIsign v3.0 changes parameters, integer representation, lattice reduction, ideal representation, ideal-to-isogeny conversion, and response sampling [5]. Its official source provides complete Cortex-M4 KeyGen, Sign, and Verify paths[20]. The v2 `find_uv` candidate table is absent from the v3 path, so the norm sketch has no direct v3 target. Large, sequentially used local states remain, which motivates the separate lifetime experiment in section 6.

3.2 Why fixed-capacity integers

GMP and its smaller mini-GMP implementation represent integers with separately allocated limb arrays and grow those arrays through an allocator [25]. Mini-GMP passed the functional response checks in a separate host experiment, but a simple 16-byte-aligned static-pool allocator needed 317,696 bytes for the measured 100-vector signing corpus and saw up to 4,830 live blocks. This host threshold is neither a mini-GMP lower bound nor an RP2350 total. We choose fixed-capacity arrays because their ownership, alignment, and reservation can be audited before linking, not because fixed capacity implies constant-time arithmetic. Used-limb scans, Euclid loops, comparisons, and exponentiation remain variable-time.

3.3 Target and memory quantities

The target is a Raspberry Pi Pico 2 with one Arm Cortex-M33 core of its RP2350, clocked at 150 MHz[26]. Code and read-only data execute from internal flash; only the chip’s 532,480 bytes of SRAM are counted as available data memory. The builds use Pico SDK 2.3.0 and Arm GNU Toolchain 15.2.1 in Release mode.

We distinguish five quantities: operation-arena payload, guarded owner, linked mutable regions, reserved Process/Main Stack Pointer regions, and an observed stack watermark. A watermark is the overwritten extent for one execution, not a bound. “Unreserved SRAM” is computed after subtracting the full PSP and MSP reservations, never only the observed watermarks.

3.4 Security boundary

Functional failures include output disagreement, boundary-canary corruption, uncleared workspace state, or an unintended allocator/GMP dependency. For side channels, we consider an observer of time, control flow, addresses, power, or electromagnetic emanations. This study records host timing and software traces and selected RP2350 timing diagnostics. It records no full-Sign power or EM trace. Constant-time execution, masking, fault resistance, production randomness, and resistance to key recovery are not claimed.

4 Lifetime Scheduling Model

Let $O = \{o_1, \dots, o_n\}$ be the operation objects and $\mathcal{T}$ the set of allowed traces, including retry and failure returns. Object o_i has size s_i , alignment a_i , type τ_i , and a secrecy attribute. Its live points on trace t are $L_i(t)$. We define the interference graph $\mathcal{G} = (O, E)$ by

$$(o_i, o_j) \in E \iff \exists t \in \mathcal{T} : L_i(t) \cap L_j(t) \neq \emptyset. \quad (1)$$

Thus, disjointness on one successful execution is insufficient.

For offsets $x_i \geq 0$, a valid typed placement satisfies

$$x_i \equiv 0 \pmod{a_i}, \quad (2)$$

$$(o_i, o_j) \in E \implies (x_i + s_i \leq x_j) \vee (x_j + s_j \leq x_i). \quad (3)$$

The arena height is $M_{\text{arena}} = \max_i(x_i + s_i)$. We do not solve the general optimization problem. Reviewed phase annotations yield a feasible first-fit placement, emitted as typed C unions/structures with `sizeof`, `offsetof`, alignment, and static-assert checks.


For pre/post guards $g_{\text{pre}}, g_{\text{post}}$, PSP and MSP reservations $R_{\text{PSP}}, R_{\text{MSP}}$, and other exclusive mutable state R_{other} , the firmware quantity is

$$R_{\text{excl}} = (M_{\text{arena}} + g_{\text{pre}} + g_{\text{post}}) + R_{\text{PSP}} + R_{\text{MSP}} + R_{\text{other}}. \quad (4)$$

For v2 these terms are 172,080, 128, 131,072, 8,192, and 9,936 bytes, giving 321,408 bytes.

Two objects may share an offset only when their reviewed contracts establish: (i) disjoint lifetimes on success, retry, and failure paths; (ii) no pointer escape or returned reference; (iii) required clearing before reuse; (iv) typed aliasing and alignment; and (v) cleanup on every return edge. ABI probes, `-Wvla` `-Walloca` `-Werror`, ASan/UBSan, canaries, zeroization fixtures, and ELF audits serve as independent failure gates. They do not constitute whole-program alias verification.

(a) Logical lifetimes



(b) Physical layout



Figure 2: v2 logical lifetimes and their typed physical layout. The 77,168-byte lattice state overlaps none of the `find_uv` phases. Three mutually exclusive phase members share the 94,912-byte phase union.

KeyGen, Sign, and Verify also share one operation owner because the embedded API serializes them. This makes the reservation their maximum rather than their sum. Reentrancy, concurrent signing, and cross-core use are outside the API contract.

5 The v2 Method

5.1 Reducing resident state

The starting `find_uv` stores candidate vectors, full-precision norms, quotients, row counts, and sorting records simultaneously. We apply the following transformations, checking differential transcripts after each stage:

1. pack each four-coordinate candidate in `int8_t[4]` after verifying that each coordinate lies in $[-2, 2]$;
2. sort a `uint16_t` index permutation rather than candidate copies;
3. reconstruct per-candidate quotients at their use sites;
4. stream two rows instead of retaining all seven;
5. move large variable-length locals into caller-owned typed workspaces;
6. overlay ML2, candidate, fixed-degree, theta, discrete-log, inversion, and integer-temporary phases where the lifetime contracts permit it;

Table 1: v2 ablation. Rows before explicit lattice ownership describe the candidate region; later rows describe caller-owned workspace or the operation arena, and are not whole-firmware peaks.

Configuration	Region [B]	Change introduced
Compact starting point	6,339,868	Five simultaneous heap payloads
Packed coordinates	2,044,252	Four-byte candidates and packed records
Index sorting	1,905,728	Reusable <code>uint16_t</code> permutation
Deferred quotient	962,240	Reconstruct at use sites
Two-row streaming	275,840	Retain two of seven rows
Explicit lattice owner	353,008	Move a 77,168-byte VLA into the owner
Full-norm-table baseline	353,008	Overlay other phases; no dynamic Arm frame
v2 proposed	172,080	Sketch rows plus full-precision fallback

7. share the serialized KeyGen/Sign/Verify owner; and

8. replace the two full-precision norm rows with compact sketches and a full-precision fallback.

Moving the lattice VLA to an owner intentionally increases the displayed region from 275,840 to 353,008 bytes: this is an ownership checkpoint, not a claimed saving. It prevents a change of storage class from being misreported as reduced total memory.

5.2 Order-preserving norm sketch

For a candidate v , the positive integer norm $N(v)$ can be reconstructed from its retained packed coordinates, the row Gram matrix G_j , and adjusted norm d_j . Let

$$\ell(n) = \begin{cases} 0, & n = 0, \\ \lfloor \log_2 n \rfloor + 1, & n > 0, \end{cases} \quad (5)$$

and, for $w = 64$,

$$h_w(n) = \begin{cases} n, & \ell(n) \leq w, \\ \lfloor n/2^{\ell(n)-w} \rfloor, & \ell(n) > w. \end{cases} \quad (6)$$

The resident sketch is $S_w(n) = (\ell(n), h_w(n))$. Candidate comparison uses bit length, the retained 64-bit prefix, reconstructed full-precision norms on a sketch tie, and the original enumeration index, in that order.

Proposition 1 (Order preservation). *Suppose $N(v_i)$ can be reconstructed with full precision for every retained candidate. The sketch comparator orders candidates as the lexicographic key $(N(v_i), i)$ does.*

Proof. Integers with different bit lengths have the same order as their lengths. For equal lengths and unequal 64-bit prefixes, the most significant differing bit lies in a retained position, so the prefix decides the integer order. If both sketch fields agree, the comparator reconstructs and compares the two full-precision integers. It uses the enumeration index only when those integers are equal. These cases cover every pair. $\square$

For example, 2^{100} and $2^{100} + 1$ have equal sketches but are sent to the reconstruction path. Correctness therefore does not depend on assuming that sketch ties are rare. In the attribution corpus, reconstruction was used for 5,041 of 2,062,361 sorting comparisons (0.2444%); this frequency is a performance observation only.

Algorithm 1 gives a compact, self-contained view of the two central changes. `PREPAREROW` generates candidates in the reference order and sorts their index permutation with `COMPARE`. `SEARCHPAIRS` retains the reference row-pair, in-row, and congruence orders. The detailed source obligations and failure cases appear in table 2.

Algorithm 1 Order-preserving comparison and two-row `find_uv` schedule


Require: packed candidates, row descriptors $R_j = (G_j, d_j)$, target T , and caller-owned workspace W

Ensure: reference-key comparison; FINDUV-LM clears W on every exit

```
1: function COMPARE( $i, j, V, S, G, d$ )
2:   if  $i = j$  then
3:     return 0
4:   else if  $S[i].\ell \neq S[j].\ell$  then
5:     return CMP( $S[i].\ell, S[j].\ell$ )
6:   else if  $S[i].h \neq S[j].h$  then
7:     return CMP( $S[i].h, S[j].h$ )
8:   end if
9:    $n_i \leftarrow \text{FULLNORM}(V[i], G, d); n_j \leftarrow \text{FULLNORM}(V[j], G, d)$ 
10:  if  $n_i = \text{FATAL}$  or  $n_j = \text{FATAL}$  then
11:    return FATAL
12:  else if  $n_i \neq n_j$  then
13:    return CMP( $n_i, n_j$ )
14:  else
15:    return CMP( $i, j$ )
16:  end if
17: end function
18: function FINDUV-LM( $T, \mathcal{I}, W$ )
19:   $(s, L, R_0, \dots, R_{r-1}) \leftarrow \text{PREPARELATTICE}(T, \mathcal{I}, W)$ 
20:  if  $s = \text{FATAL}$  then
21:    return FINISH(FATAL,  $W$ )
22:  end if
23:  CLEARPHASESTORAGE( $W$ )
24:  for  $j \leftarrow 0$  to  $r - 1$  do
25:     $c[j] \leftarrow \text{PREPAREROW}(j, 0, W)$  ▷ validation pass
26:    if  $c[j] = \text{FATAL}$  then
27:      return FINISH(FATAL,  $W$ )
28:    end if
29:  end for
30:  for  $j_1 \leftarrow 0$  to  $r - 1$  do
31:     $m_1 \leftarrow \text{PREPAREROW}(j_1, 0, W)$ 
32:    if  $m_1 \neq c[j_1]$  then
33:      return FINISH(FATAL,  $W$ )
34:    end if
35:    for  $j_2 \leftarrow j_1$  to  $r - 1$  do
36:      if  $j_2 = j_1$  then
37:         $(V_2, m_2) \leftarrow (W.V[0], m_1)$  ▷ diagonal alias
38:      else
39:         $m_2 \leftarrow \text{PREPAREROW}(j_2, 1, W); V_2 \leftarrow W.V[1]$ 
40:        if  $m_2 \neq c[j_2]$  then
41:          return FINISH(FATAL,  $W$ )
42:        end if
43:      end if
44:       $z \leftarrow \text{SEARCHPAIRS}(T, W.V[0], V_2, R_{j_1}, R_{j_2})$ 
45:      if  $z = \text{FATAL}$  then
46:        return FINISH(FATAL,  $W$ )
47:      end if
48:      if  $z$  is a solution then
49:        if not VALIDATEANDPUBLISH( $z$ ) then
50:          return FINISH(FATAL,  $W$ )
51:        end if
52:        return FINISH(SUCCESS,  $W$ )
53:      end if
54:    end for
55:  end for
56:  return FINISH(NOTFOUND,  $W$ )
57: end function
```

Here FULLNORM evaluates the full-precision quadratic form, requires a positive quotient with zero remainder on division by d , and otherwise returns FATAL. FINISH clears the complete caller-owned workspace before returning.

(a) Construction



(b) Comparison




Figure 3: Resident sketch and comparison path. A tie causes reconstruction of both full-precision norms before the comparator returns an order.

Rows contain at most 624 candidates and are sorted by an in-place heapsort of their index permutation. If norm reconstruction fails, the comparison callback latches a fatal status; an index order is returned only to finish the sort, after which the entire row is rejected. The temporary order is never published as a successful result.

5.3 From the proposition to the C implementation

The frozen full-norm-table reference is commit `6b79cfb5...`. The proposed commit `71099e08...` is its direct child. A source audit binds 22 predicates and anchors to their digests. Table 2 summarizes the obligations connecting proposition 1 to the C search.

Table 2: Correspondence between the v2 full-norm-table reference and the proposed `find_uv`.

Obligation	Reference	Proposed	Failure handling
Row generation	Same coordinate order, sign representative, 2/3 filters, order-4 representative, odd norm	Same loops and predicates, followed by <code>finduv_eval_norm</code>	Nonintegral division, non-positive value, or capacity overflow is fatal
Sort key	Full-precision $(N(v), i)$	Length, top 64 bits, reconstructed norm, then i	Reconstruction error is latched and rejects the row
Row pairs	Increasing j_1 , then $j_2 \geq j_1$	Same triangular order; diagonal aliases one row	A regenerated count mismatch is fatal
In-row order	Increasing i_1 ; $i_2 = i_1$ on diagonal and 0 otherwise	Same starts and increments	Norm reconstruction error is fatal
Congruence sequence	Same initial residue, stride n_1 , and $v < \lfloor T/n_2 \rfloor$	Same inverse, bound, and loop order	Noninvertible continues; damaged arithmetic is fatal
Publication	Validate and publish only after success	Same publication structure; Success/Not-Found/Fatal remain distinct	Common terminal cleanup clears the workspace

Proposition 2 (First accepted result). *For a fixed input, assume arithmetic succeeds in both implementations and the correspondences in table 2 hold. The proposed search returns the same first accepted tuple as the reference, or both return NOTFOUND within the reference search domain.*

Proof. By proposition 1, each reconstructed row has the reference candidate order. A validation pass retains only row counts, and the search pass reconstructs rows deterministically. The row-pair, in-row, and congruence orders are the same by table 2; hence both searches first accept the same tuple or exhaust the same sequence. $\square$

This result is conditional on the frozen source correspondence. It is not a formal proof of all C behavior. The source audit is therefore combined with 12 reference/proposed transcripts, the two 100-request differential runs, an intentional sketch-tie fixture, and NotFound, retry, fatal, guard, and clearing tests.

5.4 Lifetime certificate and linked closure

The final v2 firmware fixes 52 project translation units reachable from the Level-I/RADIX32 operations. ELF checks reject allocator, GMP, system-RNG, legacy large-stack, unresolved-symbol, or out-of-region sections; the linked heap is zero bytes. Eleven principal regions have a machine-readable lifetime certificate recording

type, Arm size/alignment, offset, construction and last-use anchors, prohibited co-residents, pointer passing and escape review, cleanup on success/failure/retry, and associated static/dynamic tests.

Table 3: Summary of principal v2 lifetime-certificate records.

Object	Size [B]	Align [B]	Exclusion contract	Exit checks
Operation owner	172,080	8	Serialized K/S/V	Full zero scan and guards
KeyGen union	172,080	8	<code>find_uv</code> versus discrete log	Terminal clear, retry contract, assertions
Sign union	172,080	8	Seven principal phase members	Terminal clear, guards, 100-request tests
<code>find_uv</code> lattice	77,168	8	Coexists with candidates; no overlay	Offset assertions and ASan/UBSan
<code>find_uv</code> phase union	94,912	8	Lattice multiply, candidate, fixed degree	One-way transitions and all terminal exits
Verify union	15,428	4	Even-isogeny versus theta	Wrapper cleanup and modified-signature fixture

The pointer-escape scan rules out direct global/static stores in the frozen source, but is not a sound whole-program alias analysis. In particular, an outer Sign retry clears/finalizes the active callee and relies on subsequent initialization; a full-owner zero scan is performed at the public operation’s terminal exit. The certificate records this distinction rather than silently treating every retry as a whole-owner clear.

6 Transfer to SQIsign v3

The official v3 `p324_3/m4f` object reports a 34,816-byte frame for `quat_111_dual_reduce_ideal`. D1, an intermediate one-function variant, introduces two typed overlays:

1. lattice-reduction state `fp_ldl_t` shares storage with the integer matrix `Aout`, which is first constructed after reduction; and
2. `Ainv_total`, after deriving `Aout`, shares its equal sized region with the later output basis `Bout`. The still-live Gram matrix and transform do not overlap it.

D3 adds a third overlay in `protocols_sign`: the commitment ideal’s last reference precedes the first write of the challenge-isogeny output product. Source-anchored certificates check those boundaries and the equal type sizes. D3 changes ownership and placement in two functions; it does not change the API, algorithms, branch conditions, or randomness consumption. Its source commit is `874658c6...`, based on official v3 commit `6d017708...`. The RP2350 firmware was built from clean harness commit `5df7acfb...`; source tree, patch, bundle, generated tree, ELF, UF2, map, archive, and capture digests are stored together.

This transfer remains local. It does not establish an automatic whole-v3 placement, and the D3 image retains 19 compiler-reported dynamic frames. A separate D2 fixed-frame prototype is used only to study stack analysis; its results do not transfer to the D3 binary.

7 Evaluation Method

7.1 Version isolation

v2, official v3, D1, D2, and D3 use separate source, build, and result directories. Within-v3 comparisons share the official base, parameter set, toolchain, and options. Cross-generation timing or PSP differences are descriptive only because parameters, algorithms, integer backends, and memory ownership change together. In particular, the explicit v2 arena and the v3 PSP watermark are not subtracted as though they were the same resource.

7.2 Functional gates

For v2, 12 frozen transcripts compare the full-norm-table and proposed implementations. One ordinary-API process and two separately started low-memory processes each consume all 100 official v2 Level-I requests. Each vector checks public/secret keys, signed message, post-RNG state, successful Open, one-bit modified-signature rejection, guards, and terminal clearing. All three generated response streams are byte-identical. The expected response is produced by the ordinary path at the same compact-source commit. The two low-memory runs therefore test fresh-process repeatability, but do not provide independent oracles. We call this a differential-conformance test using official requests, not a historical-response KAT.

For v3, official and D3 builds each pass the NIST API test, scheme self-test, and 100 official-response vectors for p324_3, p500_27, and p664_17. This is 300 distinct parameter-set/request pairs run once on each implementation, or 600 implementation-vector runs in total. D1 additionally has five interleaved official/adapted target pairs and a 10-vector campaign in two linked-code placements. D3’s final target evidence is one clean p324_3 vector-0 K/S/V capture. These corpora are finite regression gates, not all-input equivalence proofs.

7.3 Stack method

The v2 stack analysis uses 126 GCC stack-usage files with 1154 records and no dynamic record. It resolves every linked, input-dependent Thread-mode path from the three operation roots, including indirect calls and manually bound assembly/runtime frames. Two recursive components halve their length at each child; with maximum length 248 they have at most nine active frames. A source, constant, object, or function-body digest change stops the audit.

The longest condensed-call-graph path receives a conservative 212-byte allowance for one maximum Secure Armv8-M exception entry [27]. The result concerns the operation PSP. Handler software uses the MSP and is not included; neither are the enabled-interrupt set, nesting policy, or live MSP at entry.

The separate v3 D2 analysis follows the same pattern after replacing 19 dynamic frames with fixed-capacity frames. It closes the linked operation roots, but retains 18 indirect callback sites in candidate interrupt-handler paths. Hence neither the v2 nor D2 result is a whole-program, interrupt-inclusive stack bound.

8 Results

8.1 v2 SRAM and flash

The two full-precision norm rows occupy 269,568 bytes. Removing them reduces the candidate workspace from 275,840 to 14,024 bytes. The phase union stops at 94,912 bytes because an ML2 retry workspace becomes the largest member. Together with the 77,168-byte long-lived lattice state, the complete operation arena is therefore

$$77,168 + \max(94,912, 14,024) = 172,080 \text{ bytes.} \quad (7)$$

The full-norm-table baseline is 353,008 bytes, so the reduction is 180,928 bytes, or 51.2532%.

Table 4: Linked v2 proposed firmware memory accounting.

Quantity	Bytes	Scope
On-chip SRAM total	532,480	main + scratch
Operation-arena payload	172,080	KeyGen/Sign/Verify union
Guarded operation owner	172,208	main bank
PSP reservation	131,072	main bank
Other main-bank state	9,936	data and runtime
Main-bank linked use	313,216	preceding rows
MSP reservation	8,192	scratch bank
Exclusive SRAM reservation	321,408	−180,928 B
Unreserved SRAM	211,072	+180,928 B
Heap	0	no allocator symbol
BIN image	288,384	+880 B

The arena reduction propagates one-for-one to exclusive SRAM and unreserved SRAM; it is not a move from heap to stack. The proposed BIN image is 288,384 bytes, 880 bytes larger than the baseline, reflecting the reconstruction and comparison code.

8.2 v2 operation PSP and time

Two boots of the same stored UF2 and deterministic input observe PSP extents of 91,980, 120,452, and 20,768 bytes for KeyGen, Sign, and Verify. They also reproduce the transcripts and stack extents. This is a cross-check, not input-population sampling.

The smallest PSP margin is 10,408 bytes for Sign. The analysis covers all linked input-dependent Thread-mode paths in this image plus one maximum exception entry, but excludes handler frames, nesting, and MSP usage.

Table 5: v2 operation-PSP analysis for the fixed image. The final column is relative to the 131,072-byte reservation.

Operation	Observed [B]	Software path [B]	PSP result [B]	Margin [B]
KeyGen	91,980	91,980	92,192	38,880
Sign	120,452	120,452	120,664	10,408
Verify	20,768	20,768	20,980	110,092

Table 6: First RP2350 boot for the deterministic v2 operations. Cycles are elapsed microseconds multiplied by 150, not hardware-counter readings.

Operation	Time [s]	Cycles at 150 MHz	PSP extent [B]
KeyGen	2,698.150029	404,722,504,350	91,980
Sign	7,611.258527	1,141,688,779,050	120,452
Verify	0.814341	122,151,150	20,768

KeyGen takes about 45 minutes and Sign about 127 minutes, so the v2 result is a memory-feasibility result rather than a real-time signer. On the host, two reversed-order 30-row campaigns give proposed/baseline median ratios of 1.003695 and 1.006464 for KeyGen, and 0.999481 and 0.998980 for Sign. A small KeyGen cost is visible; a Sign slowdown direction is not reproduced. Single target observations have ratios 1.000704, 1.037312, and 1.000593, but cannot estimate a target timing distribution. The design is therefore a RAM-oriented Pareto point with an 880-byte flash cost, not a claim of dominance on every metric.

8.3 v3 local transfer

In D1, the two overlays reduce the `quat_111_dual_reduce_ideal` frame by 3,928 bytes; five paired target runs show the same reduction in Sign PSP. D3 adds the second function and produces the results in table 7.

Table 7: Official v3 versus D3 on p324_3. Frame sizes are Arm object values; official PSP is the earlier five-run median and D3 PSP is one clean observation.

Quantity	Official v3	D3	Difference
<code>quat_111_dual_reduce_ideal</code> frame [B]	34,816	30,888	-3,928
<code>protocols_sign</code> frame [B]	20,008	19,272	-736
KeyGen PSP extent [B]	56,972	56,972	0
Sign PSP extent [B]	101,060	96,396	-4,664 (-4.6151%)
Verify PSP extent [B]	37,444	37,444	0

The D3 Sign PSP watermark is 96,396 bytes, a 4,664-byte (4.6151%) reduction from official v3. The incremental 736-byte reduction from D1 to D3 equals the `protocols_sign` frame reduction. D3 has only one target capture, so its elapsed time is archived but not used for a slowdown claim.

All six function-frame comparisons decrease, while all six host implementation/parameter combinations pass their 100 official-response vectors. This supports applicability of the three local overlays across the official parameter sets. It does not supply whole-image SRAM or target watermarks for p500_27 or p664_17.

The D1 10-vector/two-placement campaign passes 40 valid K/S/V trials and 40 modified-signature rejections. PSP extents agree across both code placements, and the 3,928-byte D1 Sign reduction appears in all 20 observations. It addresses a one-vector/one-address alternative explanation for D1, but it is not a D3 multi-input campaign.

8.4 What the separate fixed-frame prototype establishes

D2 converts the 19 dynamic-frame reports to fixed-capacity frames, rejects capacity violations with fail-stop traps, and builds with VLA/alloca errors. The linked K/S/V operation-PSP results, after adding the 212-byte exception entry, are 108,300, 127,932, and 40,468 bytes. A clean vector-0 target cross-check observes 62,096, 101,060, and 40,252 bytes, respectively.

This is a negative result for the simple hypothesis that replacing VLAs with fixed arrays automatically lowers measured stack: D2 leaves the Sign watermark unchanged and increases KeyGen and Verify. More importantly, D2 is a different binary from D3. Its operation-PSP analysis says nothing about D3’s 19 dynamic frames, handler callbacks, IRQ nesting, or MSP use.

Table 8: D3 Arm frame changes for all official v3 parameter sets.

Parameter	Function	Official [B]	D3 [B]	Reduction
p324_3	quat_l1l_dual_reduce_ideal	34,816	30,888	−3,928 (11.2822%)
p324_3	protocols_sign	20,008	19,272	−736 (3.6785%)
p500_27	quat_l1l_dual_reduce_ideal	46,088	40,904	−5,184 (11.2480%)
p500_27	protocols_sign	27,304	26,328	−976 (3.5746%)
p664_17	quat_l1l_dual_reduce_ideal	63,016	55,904	−7,112 (11.2860%)
p664_17	protocols_sign	36,728	35,392	−1,336 (3.6376%)

Table 9: Stack-claim decomposition. “Established” refers only to the frozen linked image and the stated layer.

Layer	Status	Boundary
v2 K/S/V operation PSP	Established	Linked input-dependent Thread-mode paths plus one maximum PSP exception entry
v3 D2 K/S/V operation PSP	Established	Separate fixed-frame binary, linked software paths plus the same entry allowance
v3 D3 operation PSP	Open	Two local frame changes and one p324_3 watermark; 19 dynamic-frame records remain
Handler software, MSP, nesting	Open	Requires enabled-IRQ policy, callback resolution, handler frames, nesting, and entry-time MSP state

9 Side-Channel Findings

A fixed arena removes allocator-address variation but does not make loop counts, branches, or memory accesses independent of secrets. The v2 path still contains sketch-tie reconstruction, variable candidate search, rejection, used-limb scans, early-return comparison, Euclid loops, and secret-derived modular exponents.

For v2, two reversed-order fixed-versus-random host timing screens, 64 pairs each, give first-order statistics 3.570 and 3.576 and centered second-order statistics −5.410 and −5.374. Both second-order values cross the commonly used $|t| > 4.5$ exploratory threshold[28]. This small screen detects a repeatable distributional difference; a value below the threshold would not certify nonleakage. Full-Sign software traces also show reproducible edge and effective-address differences when signing randomness changes.

The known SPA route targets secret-derived exponents during Cornacchia processing[24]. In our 12-seed v2 corpus, Sign reaches the Cornacchia modular-square-root path 361 times, 190 in the modulus class that directly matches the published exponent relation. An isolated RP2350 powmod diagnostic has correlation 0.999860 between exponent Hamming weight and time, with an approximate slope of 4.004 ms per set bit. This is a timing diagnostic, not an analog trace or a full-Sign key-recovery experiment.

For v3, a fixed-message, fixed-signing-randomness, fixed-key-address screen uses ten official keys in two key orders. Across official and D1 host builds, 1,200 verified signatures reproduce key-associated median rankings; the minimum between-order Spearman coefficient is 0.987879. The corresponding RP2350 campaign collects 200 verified signatures with cache normalization; both implementations have rank coefficient 1.000000 and key-median spans of 50.8419–51.3752%. Public and secret key components change together, so the result is key-associated timing, not secret-only attribution.

Fixed-schedule exponentiation and sampling prototypes remove selected local variations but leave value-dependent or upper-level work. No prototype meets a deployment criterion. The artifact therefore records a negative side-channel-resistance decision; memory reduction and response-vector agreement do not alter it.

10 Contract Compiler and Automation Boundary

The artifact includes a small contract compiler. Each manually annotated object provides phases, size, alignment, secrecy, pointer-escape status, and cleanup events. The tool constructs the interference graph, assigns aligned offsets by first fit, and emits JSON, Graphviz, typed C unions/structures, and static assertions. On the frozen v2 annotations it reproduces the 172,080-byte layout. The generated header is compiled against the Arm types and checked against the hand-reviewed workspace offsets.

A separate fail-closed coverage gate enumerates all 4 declared physical members and 18 direct member/whole-phase access forms. Adding a member or direct access without an annotation causes failure. This closes one practical annotation-omission class.

The tools do not infer phases or lifetimes from C. They do not establish the absence of indirect aliases, pointer escape, missing failure-edge cleanup, function-pointer targets, or changes introduced by whole-

Available evidence




Figure 4: Side-channel evidence and boundary. The available experiments detect variable behavior. Full-Sign analog acquisition and target-trace key recovery remain unperformed.

program optimization. Placement soundness therefore remains conditional on manual annotation completeness. The appropriate description is a contract compiler with a coverage gate, not a general automatic placement or static alias-analysis system.

11 Limitations and Next Steps

The evidence supports a narrower conclusion than an unconstrained “embedded SQISign” claim:

1. **Target scope.** All physical measurements use one RP2350 board and one Cortex-M33 core. v2 covers Level I/RADIX32. D3 host response checks and Arm frames cover all three v3 parameter sets, but target whole-image evidence covers only p324_3.
2. **Input scope.** The v2 target result repeats one deterministic K/S/V path twice. Its operation-PSP analysis closes linked input-dependent Thread-mode paths for the fixed image, but not handler or MSP behavior. D3 has one target vector and no target timing distribution.
3. **Functional scope.** v2 differential tests and v3 official response checks are finite corpora. The C correspondence audit and proposition 2 do not form an all-input semantics proof.
4. **Manual contracts.** The coverage gate catches declared-member and direct-access omissions, but phase meaning, indirect aliases, escape, and failure cleanup remain manually reviewed assumptions.
5. **Stack scope.** A whole-program, interrupt-inclusive maximum requires resolution of handler callbacks, enabled IRQs and priorities, nesting, and live MSP state. D3 additionally requires bounds for its remaining dynamic frames.
6. **Security scope.** The implementations are variable-time. Full-Sign power/EM traces, independent acquisitions, positive-control attack reproduction, masking, and a complete constant-time rewrite remain future work.
7. **Practicality.** v2 KeyGen and Sign are too slow for many products. The 211,072 unreserved bytes may accommodate system integration or countermeasures, but that capacity has not been demonstrated with production RNG, communications, interrupts, and countermeasure state active together.

Locally executable bounded campaigns listed in the earlier research plan are complete in the artifact. The remaining work is not a hidden checklist item: it requires either a stronger program analysis (whole-program interrupt and alias/lifetime coverage), additional target campaigns (D3 input diversity and the two larger parameter sets), or external power/EM acquisition equipment. Other PQC schemes and additional microcontrollers are outside this revision and are not counted as evidence of generality.

12 Conclusion

Typed lifetime scheduling and an order-preserving norm sketch reduce the SQIsign v2 Level-I operation arena from 353,008 to 172,080 bytes. The resulting RP2350 firmware uses no GMP, heap, or external PSRAM, reserves 321,408 bytes of SRAM, and leaves 211,072 bytes unreserved. Two fresh-process repetitions of the official-request differential campaign, lifetime and source certificates, linked-image audits, and fixed-image operation-PSP analysis support this bounded result.

In SQIsign v3, three local overlays across two functions reduce all six tested Arm frames over the three official parameter sets. One clean p324_3 target path reduces the Sign PSP watermark from 101,060 to 96,396 bytes. This transfer supports the reuse of the lifetime-scheduling workflow, while the v2 norm-sketch transformation remains version-specific.

The strongest remaining qualifications are explicit: D3 lacks a linked operation-PSP bound and broad target sampling; handler/MSP/nesting analysis is open; automatic placement depends on manual annotations; and the code remains variable-time with no analog resistance evidence. The contribution is thus an auditable memory-engineering result with defined evidence boundaries, not a claim of general optimality or hardened deployment readiness.

Artifact Availability

Source-reconstruction patches, scripts, machine-readable certificates, firmware digests, and RP2350 captures are available at <https://github.com/quantumshiro/tinysqisign>.

References

- [1] G. Alagic, M. Bros, P. Ciadoux, Q. Dang, T. H. Dang, J. Kelsey, J. Lichtinger, Y.-K. Liu, C. Miller, D. Moody, R. Peralta, R. Perlner, A. Robinson, H. Silberg, D. Smith-Tone, and N. Waller. *Status Report on the Second Round of the Additional Digital Signature Schemes for the NIST Post-Quantum Cryptography Standardization Process*. Tech. rep. NIST IR 8610. National Institute of Standards and Technology, May 2026. DOI: 10.6028/NIST.IR.8610. URL: <https://doi.org/10.6028/NIST.IR.8610>.
- [2] National Institute of Standards and Technology. *Round 3 Additional Digital Signature Schemes*. 2026. URL: <https://csrc.nist.gov/projects/pqc-dig-sig/round-3-additional-signatures> (visited on 09/04/2026).
- [3] L. De Feo, D. Kohel, A. Leroux, C. Petit, and B. Wesolowski. “SQIsign: Compact Post-Quantum Signatures from Quaternions and Isogenies”. In: *Advances in Cryptology – ASIACRYPT 2020*. Vol. 12491. Lecture Notes in Computer Science. Springer, 2020, pp. 64–93. DOI: 10.1007/978-3-030-64837-4_3. URL: https://doi.org/10.1007/978-3-030-64837-4_3.
- [4] SQIsign submission team. *SQIsign: Algorithm Specifications and Supporting Documentation, Version 2.0.1*. July 2025. URL: <https://sqisign.org/spec/sqisign-20250707.pdf> (visited on 09/04/2026).
- [5] SQIsign submission team. *SQIsign: Algorithm Specifications and Supporting Documentation, Version 3.0*. Tech. rep. National Institute of Standards and Technology, Sept. 2026. URL: <https://sqisign.org/spec/sqisign-20260901.pdf> (visited on 09/04/2026).
- [6] G. J. Chaitin, M. A. Auslander, A. K. Chandra, J. Cocke, M. E. Hopkins, and P. W. Markstein. “Register Allocation via Coloring”. In: *Computer Languages* 6.1 (1981), pp. 47–57. DOI: 10.1016/0096-0551(81)90048-5.
- [7] J. Gergov. “Algorithms for Compile-Time Memory Optimization”. In: *Proceedings of the Tenth Annual ACM-SIAM Symposium on Discrete Algorithms*. ACM/SIAM, 1999, pp. 907–908. DOI: 10.5555/314500.315082.
- [8] W. Kim, J. Lee, H. Kim, and C. Lee. “SQIsign with Fixed-Precision Integer Arithmetic”. In: *Public-Key Cryptography – PKC 2026*. Vol. 16553. Lecture Notes in Computer Science. Earlier version: Cryptology ePrint Archive, Paper 2025/1649. Springer, 2026, pp. 3–30. DOI: 10.1007/978-3-032-26737-5_1. URL: https://doi.org/10.1007/978-3-032-26737-5_1.
- [9] W. Kim, C. Lee, and H. Yoo. “Compact Quaternion Algorithms for SQIsign”. In: *Cryptology ePrint Archive, Paper 2026/1031* (2026). URL: <https://eprint.iacr.org/2026/1031>.
- [10] G. Borin, M. Corte-Real Santos, J. Komada Eriksen, R. Invernizzi, M. Mula, S. Schaeffler, and F. Vercauteren. “Qlapoti: Simple and Efficient Translation of Quaternion Ideals to Isogenies”. In: *Cryptology ePrint Archive, Paper 2025/1604* (2025). URL: <https://eprint.iacr.org/2025/1604>.

- [11] Y.-F. Lai. “Qlapoti+ and More: Optimizing Isogeny-based Signatures”. In: *Cryptology ePrint Archive, Paper 2026/1640* (2026). URL: <https://eprint.iacr.org/2026/1640>.
- [12] M. Duparc, A. Leroux, and S. Schaeffler. “Qlapoty: Improved Analysis and Efficiency for Quaternionic Ideal to Isogeny Transformation”. In: *Cryptology ePrint Archive, Paper 2026/1700* (2026). URL: <https://eprint.iacr.org/2026/1700>.
- [13] M. J. Kannwischer, M. Krausz, R. Petri, and S.-Y. Yang. “pqm4: Benchmarking NIST Additional Post-Quantum Signature Schemes on Microcontrollers”. In: *Cryptology ePrint Archive, Paper 2024/112* (2024). URL: <https://eprint.iacr.org/2024/112>.
- [14] F. Carvalho Rodrigues, D. Gazzoni Filho, G. Adj, I. A. Canales-Martínez, J. Chávez-Saab, J. López, M. Scott, and F. Rodríguez-Henríquez. “Generation of Fast Finite Field Arithmetic for Cortex-M4 with ECDH and SQIsign Applications”. In: *Cryptology ePrint Archive, Paper 2025/1322* (2025). URL: <https://eprint.iacr.org/2025/1322>.
- [15] M. A. Aardal, G. Adj, A. Alblooshi, D. F. Aranha, I. A. Canales-Martínez, J. Chávez-Saab, D. L. Gazzoni Filho, K. Reijnders, and F. Rodríguez-Henríquez. “Optimized One-Dimensional SQIsign Verification on Intel and Cortex-M4”. In: *Cryptology ePrint Archive, Paper 2024/1563* (2024). URL: <https://eprint.iacr.org/2024/1563>.
- [16] L. De Feo, L.-J. Jian, T.-Y. Wang, and B.-Y. Yang. “SQIsign on ARM”. In: *Cryptology ePrint Archive, Paper 2026/394* (2026). URL: <https://eprint.iacr.org/2026/394>.
- [17] Y. Hu, S. Shen, H. Yang, and W. Wang. “Paving the Way for SQIsign: Toward Efficient Deployment on 32-bit Embedded Devices”. In: *Mathematics* 12.19 (2024), p. 3147. DOI: 10.3390/math12193147. URL: <https://doi.org/10.3390/math12193147>.
- [18] W. Wang, C. Wang, Y. Wu, Q. Xue, J. Zheng, and Y. Zhao. “Exposing SIMD Parallelism in SQIsign: An AVX-512 Implementation”. In: *arXiv preprint arXiv:2608.13948* (2026). DOI: 10.48550/arXiv.2608.13948. URL: <https://arxiv.org/abs/2608.13948>.
- [19] Anchorage Digital. *sqisign-rs: A Pure Rust Implementation of SQIsign v2.0*. Repository release 0.5.0. 2026. URL: <https://github.com/anchorageoss/sqisign-rs> (visited on 09/04/2026).
- [20] SQIsign submission team. *The SQIsign Implementation, Version 3.0*. Git tag `nist-v3`; frozen revision 6d017708... Sept. 2026. URL: <https://github.com/SQIsign/the-sqisign/tree/nist-v3> (visited on 09/04/2026).
- [21] F. Kouider, A. Mukherjee, D. Jacquemin, and P. Kutas. “Constant-time Integer Arithmetic for SQIsign”. In: *Cryptology ePrint Archive, Paper 2025/832* (2025). URL: <https://eprint.iacr.org/2025/832>.
- [22] O. Hanyecz, A. Karenin, E. Kirshanova, P. Kutas, and S. Schaeffler. “Constant Time Lattice Reduction in Dimension 4 with Application to SQIsign”. In: *Cryptology ePrint Archive, Paper 2025/027* (2025). URL: <https://eprint.iacr.org/2025/027>.
- [23] A. Basso, C. He, D. Jacquemin, F. Kouider, P. Kutas, A. Mukherjee, S. Schaeffler, and S. Sinha Roy. “Constant-time Quaternion Algorithms for SQIsign”. In: *Cryptology ePrint Archive, Paper 2025/2192* (2025). URL: <https://eprint.iacr.org/2025/2192>.
- [24] A. Mukherjee, M. Czaprynski, D. Jacquemin, P. Kutas, and S. Sinha Roy. “Simple Power Analysis Attack on SQIsign”. In: *Progress in Cryptology – AFRICACRYPT 2025*. Vol. 15651. Lecture Notes in Computer Science. Earlier version: Cryptology ePrint Archive, Paper 2025/830. Springer, 2025, pp. 245–269. DOI: 10.1007/978-3-031-97260-7_12. URL: https://doi.org/10.1007/978-3-031-97260-7_12.
- [25] T. Granlund and the GMP development team. *GNU MP: The GNU Multiple Precision Arithmetic Library*. 6.3.0. 2023. URL: <https://gmplib.org/gmp-man-6.3.0.pdf> (visited on 09/04/2026).
- [26] Raspberry Pi Ltd. *RP2350 Datasheet: A Microcontroller by Raspberry Pi*. RP-008373-DS-2. 2025. URL: <https://pip-assets.raspberrypi.com/categories/1214-rp2350/documents/RP-008373-DS-2-rp2350-datasheet.pdf?disposition=inline> (visited on 09/04/2026).
- [27] J. Yiu. *How Much Stack Memory Do I Need for My Arm Cortex-M Applications?* Arm. Mar. 2016. URL: <https://developer.arm.com/community/arm-community-blogs/b/architectures-and-processors-blog/posts/how-much-stack-memory-do-i-need-for-my-arm-cortex-m-applications> (visited on 09/04/2026).
- [28] T. Schneider and A. Moradi. “Leakage Assessment Methodology: A Clear Roadmap for Side-Channel Evaluations”. In: *Journal of Cryptographic Engineering* 6.2 (2016), pp. 85–99. DOI: 10.1007/s13389-016-0120-y. URL: <https://eprint.iacr.org/2015/207>.